

# **Towards Building An Effective Multi-Label Image Classification Model**

Jadie | Lopez | Ponce  
STINTSY S14

# Outline

- The Task & The Dataset
- Data Preprocessing
- Simple EDA
- Model Training
- Model Selection & Hyperparameter Tuning
- Data Augmentation
- Challenges w/ the Data
- Insights & Conclusion

# **The Task & The Dataset**

# The Task

- **Multi-Label Image Classification**
  - *to associate multiple **labels** with a given image*
- Classify an image according to:

*if it contains...*




tom



jerry



both



neither

# The Dataset

## *Tom and Jerry Image Classification*

- Dataset containing images from various episodes of the famous cartoon show, Tom and Jerry.
- Currently contains **5,478** images (instances)



Source: Kaggle  
<https://www.kaggle.com/datasets/balabaskar/tom-and-jerry-image-classification>

# The Dataset

## *Tom and Jerry Image Classification*

- A ground truth file (`ground_truth.csv`) is also provided containing the labels that correspond to each image file for supervised training.

```
In [120]: # read ground_truth csv file and store it as a dataframe
df_gt = pd.read_csv(ground_truth)

# display the first 5 rows only
df_gt.head()
```

|   | filename   | tom | jerry |
|---|------------|-----|-------|
| 0 | frame0.jpg | 0   | 0     |
| 1 | frame1.jpg | 0   | 0     |
| 2 | frame2.jpg | 0   | 0     |
| 3 | frame3.jpg | 0   | 0     |
| 4 | frame4.jpg | 0   | 0     |

# Data Preprocessing

# Preprocessing

- All images are in JPEG format (1280x720).
  - Resized and converted into 3D NumPy **ndarrays** (224x224x3).
    - Faster training.
    - Defined fixed input size for the models that we are going to build and train.
  - Data type for each pixel value is **float32**.
    - Values are normalized, scaling them down to 0-1 from 0-255.
    - Faster convergence.



```
In [20]: # single Tom and Jerry image (200x200x3)
         x[0]
```

```
array([[0., 0., 0.],
       [0., 0., 0.],
       [0., 0., 0.],
       ...,
       [0., 0., 0.],
       [0., 0., 0.],
       [0., 0., 0.]])
```



# Preprocessing

- With the ground truth file read and converted into a dataframe, this is used to create the ground truth labels numpy (np) array which is "y".
- Two np arrays were created: X and y.

```
In [214]: y
          array([[0, 1],
                 [0, 1],
                 [0, 1],
                 ...,
                 [1, 1],
                 [1, 1],
                 [1, 1]])
```

```
In [117]: def create_img_dataset(img_dir, ground_truth, size=[224, 224]):
           # np ndarray containing the tom and jerry images represented 3D numpy ndarrays
           X = []
           # np ndarray containing the ground truth (y) labels for each image
           y = []

           # get ground truth labels and store them into a dataframe
           labels = pd.read_csv(ground_truth)

           # get and pre-process all images from each folder in the specified image
           for sub_dir in os.listdir(img_dir):
               for img_file in os.listdir(os.path.join(img_dir, sub_dir)):
                   # get image path
                   img_path = os.path.join(img_dir, sub_dir, img_file)
                   # read image, it will be on the default BGR format first
                   img = cv2.imread(img_path)
                   # convert to RGB format
                   img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
                   # resize image according to set image dimensions
                   img = cv2.resize(img, (size[0], size[1]), interpolation=cv2.INTER_AREA)
                   # convert image to numpy array with float32 as data type
                   img = np.array(img).astype('float32')
                   # normalize the image array's values to 0-1
                   img /= 255
                   # store image in array
                   X.append(img)
                   # store image's ground truth class label, there must be a 'filename' column which
                   # corresponds to the name of the current image file being stored
                   y.append(labels[labels['filename'] == img_file].values[:, 1:].squeeze())

           # return the created image dataset as numpy ndarrays
           return np.array(X), np.array(y).astype(int)
```

```
# store image's ground truth class label, there must be a 'filename' column which
# corresponds to the name of the current image file being stored
y.append(labels[labels['filename'] == img_file].values[:, 1:].squeeze())
```

# Simple EDA

# Unique Values

- Unique values per row of `ground_truth.csv`

`[0, 0]` -> Neither

`[0, 1]` -> Jerry only

`[1, 0]` -> Tom only

`[1, 1]` -> Both

```
In [6]: # get all unique rows (1D numpy arrays) and the number of images for each
        unique, num_of_imgs = np.unique(y, axis=0, return_counts=True)

        Lets view the unique rows from the y numpy array.

In [7]: unique

array([[0, 0],
       [0, 1],
       [1, 0],
       [1, 1]])
```




tom  
[1,0]



Jerry  
[0,1]



Both  
[1,1]



Neither  
[0,0]

# Distribution of Images

- A total of **5,478** images w/ the ff. distribution:

**1,528** -> Neither

**1,240** -> Jerry


**1,930** -> Tom

**780** -> Both



# Some random samples

- Some things worth taking note of:
  - Scenes have **varying** contexts (i.e., different backgrounds, lighting, orientation)
  - In some, Tom and/or Jerry is even **physically distorted**.



# Model Training

# Splitting the Dataset

| Set            | Percentage (%) | Image Count |
|----------------|----------------|-------------|
| Training Set   | 80%            | 4,382       |
| Validation Set | 10%            | 548         |
| Test Set       | 10%            | 548         |

# Model of Choice

- **Model of Choice:** *Convolutional Neural Networks (CNNs)*
- **Specific Models Used:**
  - **AlexNet** (Krizhevsky et al, 2012)
    - Same architecture; trained from scratch
  - **ZFNet** (Zeiler & Fergus, 2014)
    - Same architecture; trained from scratch
  - **InceptionNetV3** (Szegedy et al., 2016)
    - Slightly modified architecture; with *transfer learning*



# Model Accuracies

| CNN             | Train Accuracy (%) | Test Accuracy (%) |
|-----------------|--------------------|-------------------|
| AlexNet         | 94.33%             | 88.96%            |
| ZFNet           | 93.27%             | 85.95%            |
| InceptionNetV3* | 94.37%             | 84.58%            |

\* with transfer learning

# Model Accuracies

| CNN             | Train Accuracy (%) | Test Accuracy (%) |
|-----------------|--------------------|-------------------|
| AlexNet         | 94.33%             | 88.96%            |
| ZFNet           | 93.27%             | 85.95%            |
| InceptionNetV3* | 94.37%             | 84.58%            |

\* with transfer learning

# **Model Selection and Hyperparameter Tuning**

# Hyperparameter Tuning

- **Keras Tuner** was used to conduct hyperparameter tuning for each CNN model.
- With Keras Tuner, the **random search** method was used.

## For each model:

- Did not touch and tune the original architecture (layers) themselves.
- AlexNet -> **Learning rate** of the Adam optimizer
- ZFNet -> **Learning rate** of the Adam optimizer
- InceptionNetV3 -> **Learning rate** of the Adam optimizer and **dropout rate** of the last dropout layer.

# Hyperparameter Tuning: AlexNet

```
# hp here is the parameter passed by Keras Tuner for tuning the specified hyperparameters  
def alex_net_tuner(hp):  
    alexnet = alex_net()  
  
    # tune the learning rate of the adam optimizer  
    # we use hp.Choice() to define the learning rate values to be used for each training trial  
    hp_learning_rate = hp.Choice('learning_rate', values=[1e-2, 1e-3, 1e-4])  
  
    alexnet.compile(optimizer=Adam(learning_rate=hp_learning_rate),  
                    loss='binary_crossentropy',  
                    metrics=['binary_accuracy'])  
  
    return alexnet
```

```
tuner_alexnet = kt.RandomSearch(  
    alex_net_tuner, objective='val_loss', directory='tuned_models', project_name='tuned_alexnet')
```

# Hyperparameter Tuning: ZFNet

```
def zf_net_tuner(hp):  
    zfnet = zf_net()  
  
    # tune the learning rate of the adam optimizer  
    hp_learning_rate = hp.Choice('learning_rate', values=[1e-2, 1e-3, 1e-4])  
    zfnet.compile(optimizer=Adam(learning_rate=hp_learning_rate),  
                  loss='binary_crossentropy',  
                  metrics=['binary_accuracy'])  
  
    return zfnet
```

```
tuner_zfnet = kt.RandomSearch(  
    zf_net_tuner, objective='val_loss', directory='tuned_models', project_name='tuned_zfnet')
```

# Hyperparameter Tuning: InceptionNetV3

```
def inceptionv3_tuner(hp):
    pretrained = InceptionV3(weights="imagenet", include_top=False, input_shape=(224, 224, 3))

    for layer in pretrained.layers:
        layer.trainable = False

    model = pretrained.output
    model = MaxPool2D(pool_size=(5,5), strides=(2,2))(model)
    model = Flatten()(model)
    model = Dense(4096, activation="relu")(model)

    # tune the dropout rate
    hp_dropout = hp.Float('rate', min_value=0.1, max_value=0.5, step=0.1)
    model = Dropout(hp_dropout)(model)

    output = Dense(2, activation="sigmoid")(model)

    inceptionv3 = Model(inputs=pretrained.input, outputs=output)

    # tune the learning rate for the adam optimizer
    hp_learning_rate = hp.Choice('learning_rate', values=[1e-2, 1e-3, 1e-4])

    inceptionv3.compile(optimizer=Adam(learning_rate=hp_learning_rate),
                       loss='binary_crossentropy',
                       metrics=['binary_accuracy'])

    return inceptionv3
```

```
tuner_inceptionv3 = kt.RandomSearch(
    inceptionv3_tuner, objective='val_loss', directory='tuned_models', max_trials=20,
    project_name='tuned_inceptionv3')
```

# Summary of Results: Model Selection

| Model          | Number of Correct Predictions | Test Accuracy |
|----------------|-------------------------------|---------------|
| AlexNet        | 487.5 / 548                   | 88.9599%      |
| ZFNet          | 471.0 / 548                   | 85.9489%      |
| InceptionNetV3 | 463.5 / 548                   | 84.5802%      |

| Tuned Model    | Best Optimizer Learning Rate | Best Dropout Rate | Number of Correct Predictions | Test Accuracy |
|----------------|------------------------------|-------------------|-------------------------------|---------------|
| AlexNet        | 0.0001                       | (Not tuned)       | 510.0 / 548                   | 93.0656%      |
| ZFNet          | 0.001                        | (Not tuned)       | 461.0 / 548                   | 84.1240%      |
| InceptionNetV3 | 0.0001                       | 0.5               | 481.0 / 548                   | 87.7737%      |




# Data Augmentation

# Data Augmentation

- **ImageDataGenerator** of **Keras** was used.
- The **train** set's data was augmented by:
  - **Rotating**
  - **Shifting** (Width and Height)
  - **Zooming**
  - **Flipping** (Horizontally and Vertically)

```
train_datagen = ImageDataGenerator(rotation_range=5, # rotation  
                                   width_shift_range=0.01, # horizontal shift  
                                   height_shift_range=0.01, # vertical shift  
                                   zoom_range=0.2, # zoom  
                                   horizontal_flip=True, # horizontal flip  
                                   vertical_flip=True, # vertical flip  
                                   )
```



# Data Augmentation: Results


| Model                              | Number of Correct Predictions | Accuracy |
|------------------------------------|-------------------------------|----------|
| Initially Trained AlexNet          | 487.5 / 548                   | 88.9598% |
| Tuned AlexNet                      | 510.0 / 548                   | 93.0656% |
| Tuned AlexNet w/ Data Augmentation | 518.0 / 548                   | 94.5255% |

# Challenges with the Data

# Challenge Dataset

- There are images wherein Tom and Jerry are distorted in **shape, size, and color**.
- A **challenges csv file** was added which contains the list of **some distorted images** found in the dataset.

|   | folder | image_name    |
|---|--------|---------------|
| 0 | jerry  | frame360.jpg  |
| 1 | jerry  | frame1794.jpg |
| 2 | jerry  | frame1795.jpg |
| 3 | jerry  | frame2231.jpg |
| 4 | jerry  | frame2234.jpg |



# Challenge Dataset: Results

```
In [207]: scores = best_model.evaluate(X_ch, y_ch)

1/1 [-----] - 1s 711ms/step - loss: 0.0064 - binary_accuracy: 1.0000

Let's get the best model's number of correct predictions on the challenge dataset.

In [208]: predictions = best_model.predict(X_ch)

predictions
[[ 1.00000000e+00,  9.99999999e-01],
 [ 4.1455085e-08,  9.9930751e-01],
 [ 2.7008427e-09,  9.9999905e-01],
 [ 7.0454651e-09,  9.9999970e-01],
 [ 1.9435412e-03,  9.9195880e-01],
 [ 9.5749790e-05,  9.5904613e-01],
 [ 9.9968618e-01,  3.8940880e-02],
 [ 9.9789143e-01,  2.0948276e-02],
 [ 1.0000000e+00,  5.6303469e-07],
 [ 9.9999982e-01,  1.7406048e-05],
 [ 1.0000000e+00,  6.4829583e-06],
 [ 9.9968362e-01,  1.0855265e-03],
 [ 9.9958909e-01,  1.8719690e-04],
 [ 9.9999183e-01,  1.4778913e-04],
 [ 1.0000000e+00,  3.7413497e-07],
 [ 9.9996690e-01,  7.6467968e-03],
 [ 9.9999279e-01,  6.2163710e-03],
 [ 9.9993569e-01,  1.9419657e-02],
 [ 9.9993378e-01,  2.8760826e-02],
 [ 9.9773252e-01,  7.3010395e-03],
 [ 9.9999970e-01,  7.8685865e-05],
 [ 9.9985713e-01,  2.1912361e-04],
 [ 1.0000000e+00,  1.2117609e-04],
 [ 1.0000000e+00,  2.1450575e-04]], dtype=float32)

In [209]: correct = compute_correct(predictions, y_ch)

print("Number of correct predictions: {} / {}".format(correct, len(y_ch)))
print("Test accuracy: {}".format(scores[1]))

Number of correct predictions: 32.0 / 32
Test accuracy: 1.0
```

# **Insights & Conclusion**

# Some Insights & Conclusion


**Insight #1: Hardware constraints.** Training on machines not suited for deep neural networks is a tedious task. Training time takes hours, and sometimes, some deep learning models do not train at all.

**Insight #2: Data Augmentation helps.** This greatly helps in improving the performance of a CNN model given a small dataset.



# Insights & Conclusion

**Insight #3: More challenges on the data.** Best model currently finds it difficult to correctly predict the labels on these kinds of images:



# Some Insights & Conclusion

**Insight #4:** CNNs are powerful model architectures for image **classification**. Future goal to construct and train more CNN models that are also better in performance such as **VGG-19, ResNet, DenseNet, and Efficient**.

# References:

- Baskar, B. (n.d.). Tom and Jerry Image Classification. Kaggle. Retrieved June 23, 2022, from <https://www.kaggle.com/datasets/balabaskar/tom-and-jerry-image-classification>
- Brownlee, J. (2019). A Gentle Introduction to Transfer Learning for Deep Learning. Retrieved June 22, 2022, from <https://machinelearningmastery.com/transfer-learning-for-deep-learning/>
- Draelos, R. (2019). Multi-label vs Multi-class Classification: Sigmoid vs. Softmax. Retrieved June 21, 2022, from <https://glassboxmedicine.com/2019/05/26/classification-sigmoid-vs-softmax/>
- Dwivedi, R. (2020). How Data Augmentation Impacts Performance of Image Classification, With Codes. Retrieved June 25, 2022 from <https://analyticsindiamag.com/image-data-augmentation-impacts-performance-of-image-classification-with-codes/>
- Keras (n.d.). Keras Documentation: Keras API reference. Keras. Retrieved June 26, 2022, from <https://keras.io/api/>
- Krizhevsky, A., Sutskever, I., & Hinton, G. E. (2012). ImageNet classification with deep convolutional Neural Networks. Communications of the ACM, 60(6), 84-90. <https://doi.org/10.1145/3065386>
- Lang, N. (2021). Using Convolutional Neural Network for Image Classification. Retrieved June 26, 2022, from <https://towardsdatascience.com/using-convolutional-neural-network-for-image-classification-5997bfd0ede4>
- Migmar, T. (2021). Understanding the Amazon Rainforest with Multi-Label Classification + VGG-19, Inceptionv3, AlexNet & Transfer Learning. Retrieved June 24, 2022, from <https://towardsdatascience.com/understanding-the-amazon-rainforest-with-multi-label-classification-vgg-19-inceptionv3-5084544fb655>
- Mustafeez, A. Z. (n.d.). What is early stopping? Retrieved June 23, 2022, from <https://www.educative.io/answers/what-is-early-stopping>
- Rausch, D. (2021). EDA for Image Classification. Retrieved June 24, 2022, from <https://medium.com/geekculture/eda-for-image-classification-dcada9f2567a>
- Szegedy, C., Vanhoucke, V., Ioffe, S., Shlens, J., & Wojna, Z. (2016). Rethinking the inception architecture for computer vision. In Proceedings of the IEEE conference on computer vision and pattern recognition (pp. 2818-2826).
- Zeiler, M. D., & Fergus, R. (2014). Visualizing and understanding convolutional networks. In European conference on computer vision (pp. 818-833). Springer, Cham.

# **Towards Building An Effective Multi-Label Image Classification Model**

Jadie | Lopez | Ponce  
STINTSY S14